



Technical Manual of the Ceylan-`Traces` Layer

Organisation: Copyright (C) 2008-2019 Olivier Boudeville

Contact: `about (dash) traces (at) esperide (dot) com`

Creation date: Wednesday, August 11, 2010

Lastly updated: Wednesday, February 20, 2019

Status: Work in progress

Version: 0.9.4

Dedication: Users and maintainers of the `Traces` layer.

Abstract: The role of the [Traces](#) layer (part of the [Ceylan](#) project) is to provide Erlang applications with advanced trace services, so that the user can efficiently log, browse and search through detailed runtime messages that may be emitted concurrently (i.e. in a parallel, distributed way).

We present here a short overview of these services, to introduce them to newcomers. The next level of information is to read the corresponding [source files](#), which are intensely commented and generally straightforward.

Table of Contents

Technical Manual of the Ceylan-Traces Layer	1
Overview	3
Trace Levels	3
Trace Content	4
Trace Emission	4
From member methods	5
From constructors	5
Trace Categorisation	5
Activation / Desactivation	7
Switching from basic Console Traces	7
Trace Ordering	7
Trace Output Generation	8
Trace Browsing	9
Trace Implementation	9
General Mode of Operation	9
Trace Emitters	10
Internal Trace Format	10
Licence	11
Installation of the Traces Layer	11
Support	11
Please React!	12
Ending Word	12

Overview

This layer is in charge of providing [Erlang](#) programs with the means of emitting, collecting, storing and browsing *applicative traces* (i.e. logs).

For that, various components have been designed and implemented, such as trace aggregator, emitter, listener, supervisor, etc.

They collectively constitute the [Traces](#) layer, whose prerequisites are the [WOOPER](#) layer (for object-oriented primitives) and the [Myriad](#) layer (for many lower-level services; itself a prerequisite of WOOPER).

The main purpose of this **Traces** layer is to provide adequate traces for distributed systems (a rather critical feature in order to debug in these difficult contexts), and to ease their study and browsing. A few back-ends are available for that, from the direct reading of the (raw) trace files to considerably more user-friendly solutions, such as the generation of PDF reports or the use of our more advanced trace format, which can be read by commercial tools such as [LogMX](#)¹.

Trace Levels

There are six built-in levels for trace channels, of increasing severity:

Trace Severity	Mapped Level
debug	6
trace	5
info	4
warning	3
error	2
fatal	1

There is also an addition trace severity, `void`, that designates traces that shall be muted in all cases.

Its purpose is to provide another means of muting/unmuting some traces, instead of commenting out/uncommenting said traces.

¹The Ceylan-Traces layer defined a trace format of its own, supported by our Java-based parser for LogMX.

Trace Content

The traces corresponding to an execution are represented as an wallclock-time ordered stream of trace messages.

These traces are possibly exchanged over the network or stored in a file, whose extension is conventionally `.traces`.

For example the traces for a test named `my_foobar_test` are typically stored in a `my_foobar_test.traces` file, generated by the trace aggregator in the directory from which the corresponding test was launched.

Following data is associated to a given trace:

1. **technical identifier of the emitter**, as a string (ex: `<9097.51.0>` for the PID of a distributed Erlang process)
2. **name of the emitter** (ex: `"Instance tracker"`)
3. **dotted categorization of the emitter** (ex: `"Core.Tracker.Instances"`); here for example the emitter is an element of the service in charge of the instances, which itself belongs to the tracker services, which themselves belong to the (even more general) core services
4. **application-level timestamp** (ex: operation count, relative tick, absolute timestep, or any complex, application-specific timestamp, etc.), possibly none, or undefined if not applicable (ex: a simulation that would not be started yet)
5. **wall-clock timestamp**, in the `"Year/Month/Day Hour:Minute:Second"` format (ex: `"2016/6/10 15:43:31"`); this is an emitter-side timestamp (hence not related to the wallclock time known of the trace aggregator)
6. **emitter location**, as a string (ex: the name of the Erlang node, possibly including the name of the application use case, of the user and of the host; ex: `my_foobar_test_john@hurricane.org`)
7. **dotted categorization of the trace message** itself (ex: `MyApp.MyTopic.MyTheme`)
8. **severity of the trace message** (mapped to an integer level, as discussed above)
9. the **trace message** itself, an arbitrary text of arbitrary length

Trace Emission

The following header is to be included so that an Erlang process can send traces:

```
-include("class_TraceEmitter.hrl").
```

This process can be a standalone module (ex: a test or an application launcher, see [traceManagement_test.erl](#)) or, more frequently, it might correspond to a WOOPER (active) instance, in which case it shall inherit, directly or not, from `class_TraceEmitter` (see [class_TestTraceEmitter.erl](#) for a complete example of it).

From member methods

Then sending-primitives can be used, such as:

```
?info("Hello world!")
```

or:

```
?info_fmt("The value ~B is the answer.", [MyValue])
```

Many API variations exist (see [class_TraceEmitter.hrl](#)), to account for the various [trace content](#), contexts, etc., but `?T(Message)` and `?T_fmt(MessageFormat, MessageValues)`, for `T` corresponding to a [trace severity](#), are by far the most frequently used.

From constructors

Note that for example `?debug(Message)` is a macro that expands (literally) to:

```
class_TraceEmitter:send(debug, State, Message)
```

As a result, the availability of a `State` variable in the scope of this macro is expected. Moreover, this WOOPER state variable shall be the one of a `class_TraceEmitter` instance (either directly or, more probably, through inheritance).

This is not a problem in the most common case, when using traces in member methods (as by design they should be offering such a `State`), yet in constructors the initial state (i.e. the `State` variable fed to the `construct` operator of this class) is generally not the one of a trace emitter already.

As a result, an instance will not be able to send traces until the completion of its own `class_TraceEmitter` constructor, and then it shall rely on that resulting state (for example named `TraceState`). Sending a trace from that point should be done using `?send_debug(TraceState, Message)`.

An example of some class `Foobar` inheriting directly from `TraceEmitter` will be clearer:

```
-module(class_Foobar).

construct(State, TraceEmitterName) ->
    TraceState = class_TraceEmitter:construct(State, TraceEmitterName),
    % Cannot use ?trace("Hello!"), as it would use 'State',
    % which is not a trace emitter yet! So:
    ?send_trace(TraceState, "Hello!"),
    [...]
    FinalState.
```

Trace Categorisation

In addition to browsing the produced traces per emitter, origin, theme, wallclock or applicative timestamps, etc. it is often useful to be able to sort them per **emitter categorisation**, such a categorisation allowing to encompass multiple emitter instances of multiples emitter types.

Categories are arbitrary, and are to be nested from the most general ones to the least (a bit like directories), knowing that subcategories are to be delimited by a dot character, like in: `Art.Painting.Hopper`. As a consequence, any string can account for a category, keeping in mind dots have a specific meaning.

Hierarchical categorisation allows to select more easily a scope of interest for the traces to be browsed.

For example, should birds, cats and dogs be involved, introducing following emitter categorisations might be of help:

- `Animals`
- `Animals.Birds`
- `Animals.Cats`
- `Animals.Dogs`

If wanting all traces sent by all cats to be gathered in the `Animals.Cats` trace category, one shall introduce in `class_Cat` following define *before* the aforementioned `class_TraceEmitter.hrl` include:

```
-define(trace_emitter_categorization, "Animals.Cats").
```

and use it in the constructor like the following example, where `class_Cat` inherits directly from `class_Creature`² - supposingly itself a child class of `class_TraceEmitter`:

```
-module(class_Cat).

-define(trace_emitter_categorization, "Animals.Cats").
-include("class_TraceEmitter.hrl").

construct(State, TraceEmitterName) ->
    TraceState = class_Creature:construct(State,
                                           ?trace_categorize(TraceEmitterName)),
    % Cannot use ?trace("Hello!"), as it would use 'State',
    % which is not a trace emitter yet! So:
    ?send_warning(TraceState, "Cat on the loose!"),
    [...],
    FinalState.
```

Then all traces sent by all cats will be automatically registered with this trace emitter category.

The purpose of the `trace_categorize` macro used in the above example is to register the trace categorisation define through the inheritance tree so that, from the start, the most precise category is used³.

²We chose on purpose a different classname than `class_Animal`, to better illustrate that trace categories can be freely specified.

³Otherwise, should the various constructors involved declare their own categorisation (which is the general case) and send traces, creating a cat instance would result in having these traces sorted under different emitter categories (ex: the one declared by `class_Cat`, by `class_Creature`, etc.). Tracking the messages emitted by a given instance would be made more difficult than needed.

Activation / Desactivation

The trace macros used above can be fully toggled at build-time, on a per-module basis (if disabled, they incur zero runtime overhead, and no source change is required).

See the `ENABLE_TRACES` make variable in [GNUmakevars.inc](#) for that, and do not forget to recompile all classes and modules that shall observe this newer setting.

Note that the `warning`, `error` and `fatal` trace severities will not be impacted by this setting, as they shall remain always available (never muted).

Doing so incurs a very low runtime overhead anyway (supposing of course that sending these failure-related messages happens rather infrequently), as the cost of a mostly idle trace aggregator (which is spawned in all cases) is mostly negligible - knowing that runtime resource consumption happens only when/if emitting traces for good.

Switching from basic Console Traces

In some cases, it may be convenient to have first one's lower-level, debugging traces be directly output on the console.

Then, once the most basic bugs are fixed (ex: the program is not crashing anymore), the full power of this `Traces` layer can be best used, by switching these first, basic traces to the more advanced traces presented here.

To output (basic) console traces, one may use the [trace_utils](#) module of the `Myriad` layer. For example:

```
trace_utils:debug_fmt("Hello world #~B", [2])
```

Then switching to the mainstream, more advanced traces discussed here is just a matter of replacing, for a given trace type `T` (ex: `debug`), `trace_utils:T` with `?T`, like in:

```
?debug_fmt("Hello world #~B", [2])
```

(with no further change in the trace parameters).

Trace Ordering

It should be noted that the ordering of the reported traces is the one seen by the trace aggregator, based on their receiving order by this process (not for example based on any sending order of the various emitters involved - there is hardly any distributed global time available anyway).

So, due to network and emitter latencies, it may happen (rather infrequently) that in a distributed setting a trace message associated to a cause ends up being listed, among the registered traces, *after* a trace message associated to a consequence thereof⁴; nevertheless each trace includes a wall-clock timestamp corresponding to its sending (hence expressed according to the local time of its trace emitter).

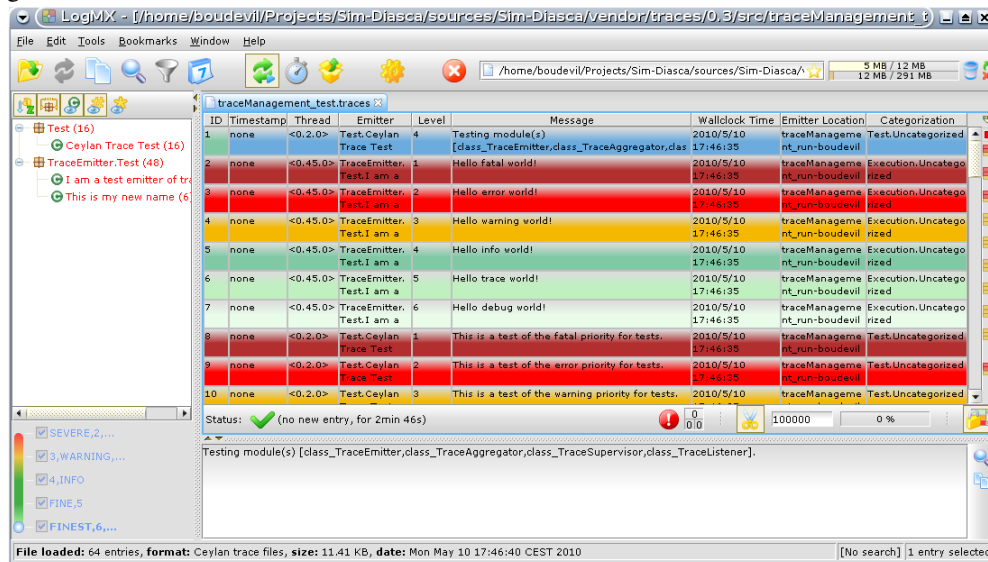
⁴A total, reproducible order on the distributed traces could be implemented, yet its runtime cost would be sufficiently high to have a far larger impact onto the executions that this trace system is to instrument than the current system (and such an impact would of course not be desirable).

Trace Output Generation

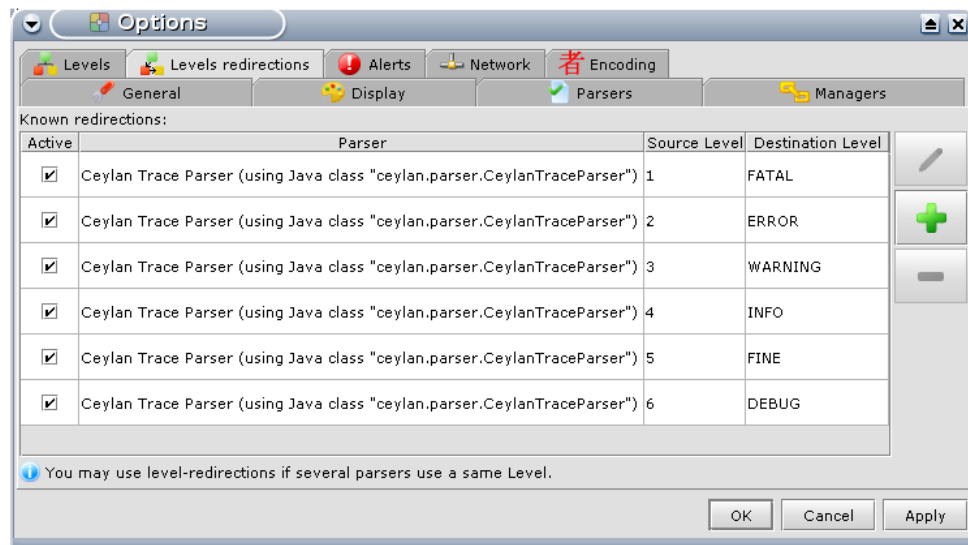
Traces may be browsed thanks to either of the following supervision solutions (see `class_TraceSupervisor.erl`):

- `text_traces`, itself available in two variations:
 - `text_only` if wanting to have traces be directly written to disk as pure, yet human-readable, text
 - `pdf`, if wanting to read finally the traces in a generated PDF file (hence the actual text includes a relevant mark-up, and as such is less readable directly)
- `advanced_traces`, for smarter log tools such as LogMX (the default), as discussed below

Indeed the most usual tool that we use for trace browsing is [LogMX](#), which we integrated:



We implemented a Java-based parser of our trace format for LogMX (see `CeylanTraceParser.java`):



Trace Browsing

Traces can be browsed with this tool:

- **live** (i.e. during the execution of the program), either from its start or upon connection to the instrumented program whilst it is already running⁵ (see `class_TraceListener.erl`)
- **post mortem** (i.e. after the program terminated for any reason, based on the trace file that it left)

The trace supervision solution can be switched at compile time (see the `TraceType` defined in `traces/src/traces.hrl`); the Traces layer shall then be rebuilt.

Trace Implementation

General Mode of Operation

All processes are able to emit traces, either by using standalone trace sending primitives (mostly for plain Erlang processes), or by inheriting from the `TraceEmitter` class, in the (general) case of **WOOPER**-based processes.

In the vast majority of cases, all these emitters send their traces to a single trace aggregator, in charge of collecting them and storing them on-disk (for most uses, their memory footprint would be quickly too large for RAM), according to an adequate trace format.

This trace format can be parsed by various trace supervisors, the most popular being **LogMX**.

Various measures have been taken in order to reduce the overhead induced by the overall trace system.

Notably traces are sent in a "fire and forget", non-blocking manner (thanks to oneways, which are not specifically acknowledged). The number of messages exchanged is thus reduced, at the cost of a lesser synchronization of the traces (i.e. there

⁵In which case the trace supervisor will first receive, transactionally, a compressed version of all past traces; then all new ones will be sent to this new listener, resulting in no trace being possibly lost.

is no strong guarantee that the traces will be ultimately recorded and displayed in the order of their emission in wallclock-time, as they will be directly and sequentially stored in their actual order of receiving by the trace aggregator⁶, an order which itself depends on the potentially varied network latencies experienced from the potential multiple sources to the trace aggregator).

Trace Emitters

When sending a trace, an emitter relies on its `trace_timestamp` attribute, and sends a (binarised) string representation thereof (obtained thanks to the `%p` quantifier of `io:format/2`). This allows the trace subsystem to support all kinds of application-specific traces (ex: integers, floats, tuples, strings, etc.).

Internal Trace Format

(for the most curious users)

Each trace line is a raw text (hence not a binary content) made of a series of predefined fields, separated by the pipe (`|`) delimiter character.

The text message included in a trace can contain any number of instances of this field delimiter.

Example of a raw trace line (end of lines added for readability):

```
<0.45.0>|I am a test emitter of traces|TraceEmitter.Test|none|  
2016/6/13 14:21:16|traceManagement_run-paul@hurricane.foobar.org|  
MyTest.SomeCategory|6|Hello debug world!
```

or:

```
<9097.51.0>|Instance tracker|Core.Tracker.Instances|14875|  
2016/6/10 15:43:31|My_application_case-john@hurricane.foobar.org|  
Execution.Uncategorized|4|Creating a new root instance tracker  
whose troubleshooting mode is enabled.
```

⁶For example, if both the trace aggregator and a process B are running on the same host, and if a process A, running on another host, emits a trace then sends a message to B so that B sends in turn a trace, then the trace from B *might* in some cases be received - and thus be listed - by the aggregator *before* the trace for A (it depends on the network congestion, relative scheduling of processes, etc.).

Licence

Ceylan-Traces is licensed by its author (Olivier Boudeville) under a disjunctive tri-license giving you the choice of one of the three following sets of free software/open source licensing terms:

- [Mozilla Public License](#) (MPL), version 1.1 or later (very close to the former [Erlang Public License](#), except aspects regarding Ericsson and/or the Swedish law)
- [GNU General Public License](#) (GPL), version 3.0 or later
- [GNU Lesser General Public License](#) (LGPL), version 3.0 or later

This allows the use of the Traces code in as wide a variety of software projects as possible, while still maintaining copyleft on this code.

Being triple-licensed means that someone (the licensee) who modifies and/or distributes it can choose which of the available sets of licence terms he is operating under.

We hope that enhancements will be back-contributed (ex: thanks to merge requests), so that everyone will be able to benefit from them.

Installation of the Traces Layer

As mentioned, the single, direct prerequisite of [Ceylan-Traces](#) is [Ceylan-WOOPER](#), which implies in turn [Ceylan-Myriad](#) and [Erlang](#), version 21.0 or more recent⁷.

Once Erlang is available, it should be just a matter of executing:

```
$ git clone https://github.com/Olivier-Boudeville/Ceylan-Myriad
$ cd Ceylan-Myriad && make all && cd ..
```

```
$ git clone https://github.com/Olivier-Boudeville/Ceylan-WOOPER
$ cd Ceylan-WOOPER && make all && cd ..
```

```
$ git clone https://github.com/Olivier-Boudeville/Ceylan-Traces
$ cd Ceylan-Traces && make all
```

Running a corresponding test just then boils down to:

```
$ cd tests && make traceManagement_run CMD_LINE_OPT="--batch"
```

Should LogMX be installed and available in the PATH, the test may simply become:

```
$ make traceManagement_run
```

Support

Bugs, questions, remarks, patches, requests for enhancements, etc. are to be reported to the [project interface](#) (typically [issues](#)) or directly at the email address mentioned at the beginning of this document.

⁷Note that, in the Ceylan-Myriad repository, we have a script to streamline the installation of Erlang, see [install-erlang.sh](#); use `install-erlang.sh --help` for guidance.

Please React!

If you have information more detailed or more recent than those presented in this document, if you noticed errors, neglects or points insufficiently discussed, drop us a line! (for that, follow the [Support](#) guidelines).

Ending Word

Have fun with Ceylan-Traces!

